



GitOps-Driven Platform Engineering:

Mastering Cluster Fleet Management



Michael Fornaro

Staff DevOps Engineer



Olga Mirensky

Lead Platform Engineer

\$ WHOAMI



Michael Fornaro

Staff DevOps Engineer

Easygo



Olga Mirensky

Lead Platform Engineer


Macquarie Group



Backstory

Stepping back...

Cloud is complicated!
Teams are swamped.



Platforms can help.

What is Platform Engineering ?!



User Experience

Consistent interface to meet where they are



Security and Compliance

Standardized and compliant by default



Golden Paths

Self-service infra using templates and automation



Management

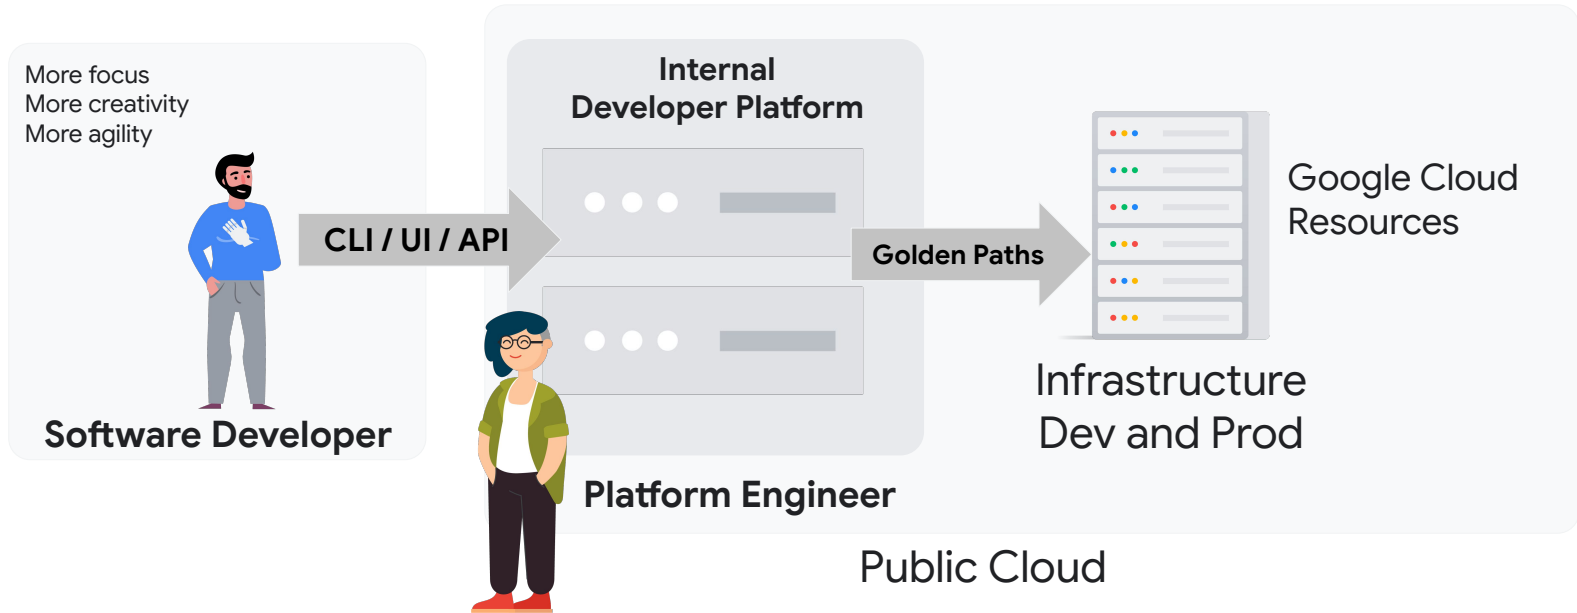
Self-service capabilities and upgrades



Platform as a Product

User-centric thinnest viable platform

Platform Engineering



Platform Engineering

More focus
More creativity
More agility




Software Developer

Define Application into
archetype



**Internal
Developer Platform**



Golden Paths



**Google Cloud
Resources**

**Infrastructure
Dev and Prod**

Platform Engineer

Public Cloud




Archetypes Introduction

Let's Go: Car Shopping

"I need a new <vehicle>!"

Really:

- A. I want to **go fast**
- B. I need to carry **lots of stuff**
- C. I don't want to **spend** much



Refinement

Once you pick a **type**, you can choose a **make** and **model**.



Customization

Then, you choose the **one for you**.




Maybe you make it even **more** for you.





NOT Car
Shopping!



Too Much Choice can be Disastrous

What we think we want

vs

How it actually turns out



https://simpsons.fandom.com/wiki/The_Homer

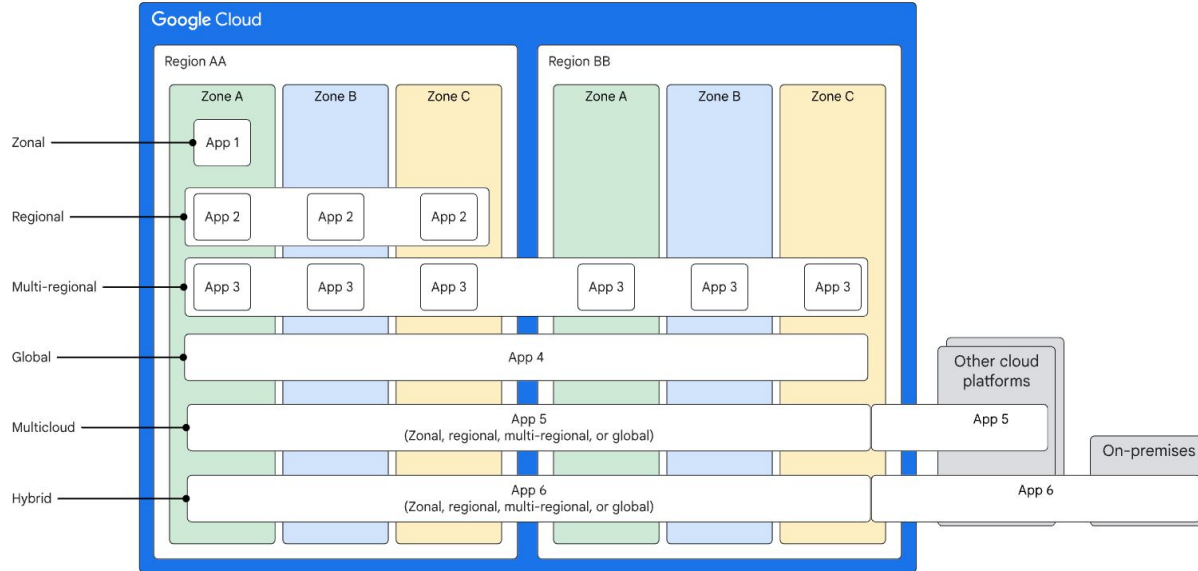
"the type of car Americans really want, not the kind we tell them they want"



A deployment archetype is an abstract, provider-independent model that you use as the foundation to build application-specific *deployment architectures* that meet your business and technical requirements


goo.gle/app-archetypes

Deployment Archetypes




goo.gle/app-archetypes

From Abstract Model to Implementation




```
achetypeyaml
1  ---
2  apiVersion: company/apps/v1alpha1
3  kind: Application
4  metadata:
5    name: whereami-frontend # Namespace, Service, Deployment, KSA, GSA, and repo names based on this.
6  spec:
7    workload:
8      application-group: "whereami" # Multiple Services belong to an application-group
9      template: "golang" # Supported values are golang, python, java, javascript
10     deployment-strategy: "global" # Strategies supported are multi-zonal, multi-regional, and global
11     # Supported regions are us-west2, us-central1 (examples)
12     # The "global" strategy ignores this field
13     locations: ["us-west2", "us-central1"]
14   service:
15     ports: ["8080", "443"]
16 > config: ...
22 > traffic: ...
25 > security: ...
28 > access: ...
33 deployment:
34   dockerfile: "/path/to/Dockerfile" # Path to the Dockerfile.
35   strategy: "canary" # Supported values are Rolling, Blue-Green and Canary, default is Rolling
36   frequency: "daily" # Supported values are Hourly, Daily, Weekly, Fortnightly, Monthly, Quarterly, default is Daily
37 infra:
38   gcp:
39     databases: ["sql", "spanner"] # Supported values are "sql", "firestore", "spanner", "redis"
40 requirements:
41   availability: "99.999%" # Target uptime percentage.
42   latency: "200ms" # Maximum acceptable response time.
43   errorRate: "0.5%" # Target maximum error rate percentage.
44   throughput: "5000qps" # Queries per second.
45   durability: "99.99%" # Data durability percentage.
46   recoveryTimeObjective: "30m" # Maximum time to recover services.
47   recoveryPointObjective: "5m" # Maximum acceptable data loss period.
48
```




How to use Archetypes?

Services can be
deployed to a
single archetype




Applications can
use services
across **multiple
archetypes**




How to use Archetypes?


Services can be
deployed to a
single archetype



Applications can
use services
across **multiple
archetypes**




Applications should
be designed for
**graceful
degradation**

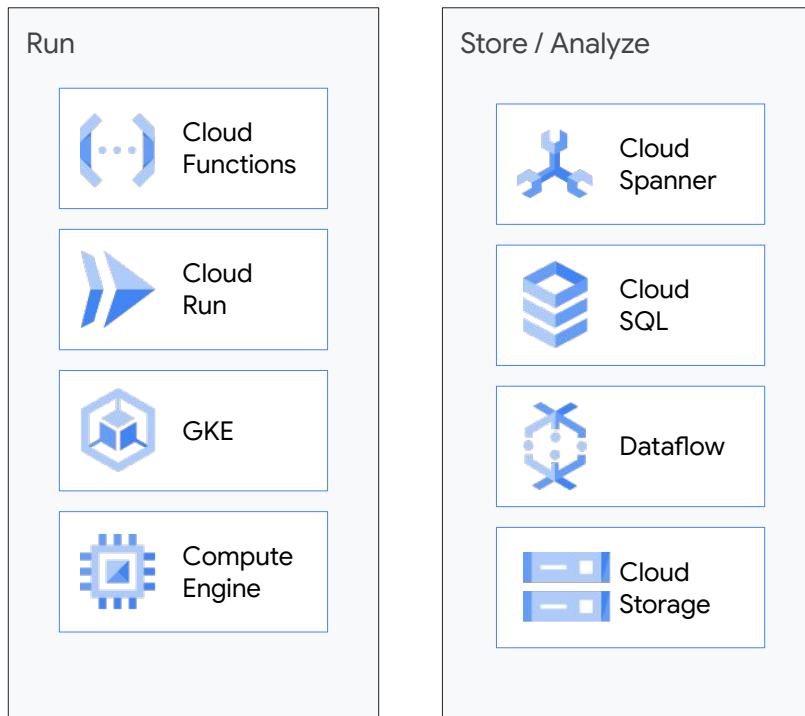


Cloud Failure Domains: **Regions**, **Zones**

Clusters, Pods, Apps, Functions, **VMs**



Products expose their failure domains differently



What is GitOps?

Declarative

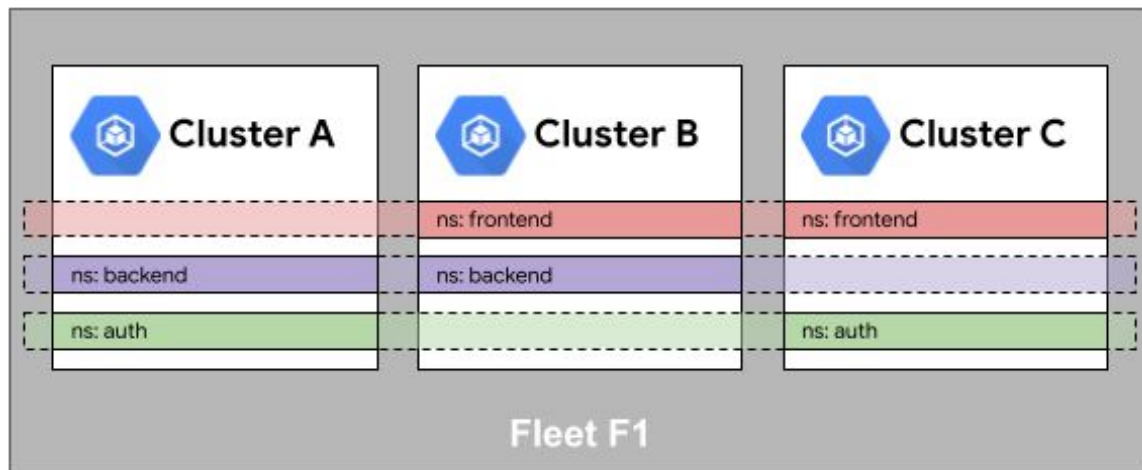
Versioned &
Immutable

Pulled
Automatically

Continuously
Reconciled

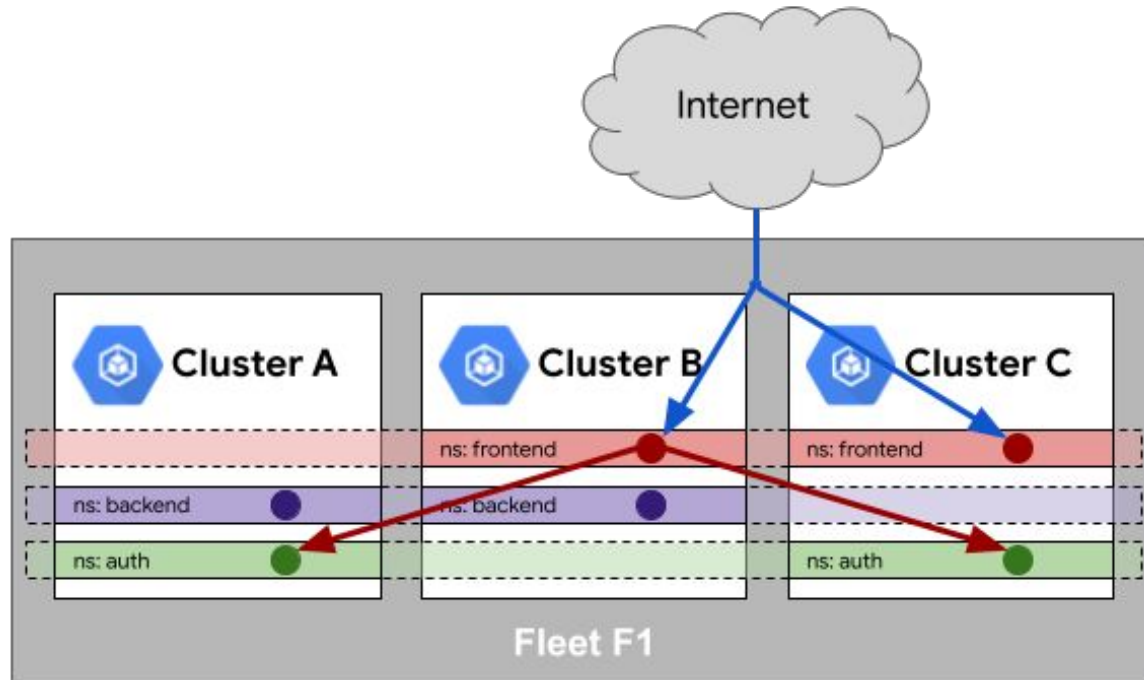
<https://opengitops.dev/>

What are Fleets of Clusters?



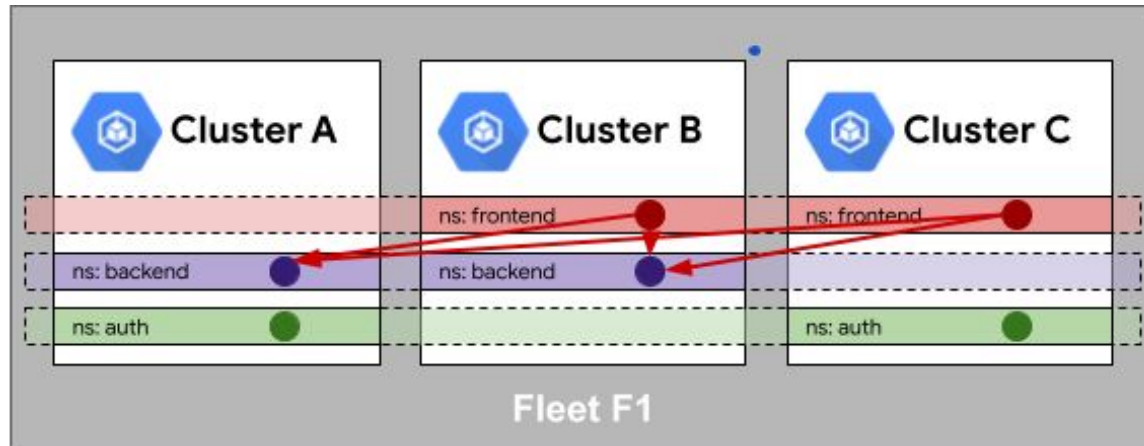
Namespace sameness in a fleet

What are Fleets of Clusters?




Service sameness in a fleet

What are Fleets of Clusters?



Identity sameness inside a fleet

Mastering Cluster Fleet Management





Demo video

Unpacking the Demo - Principles

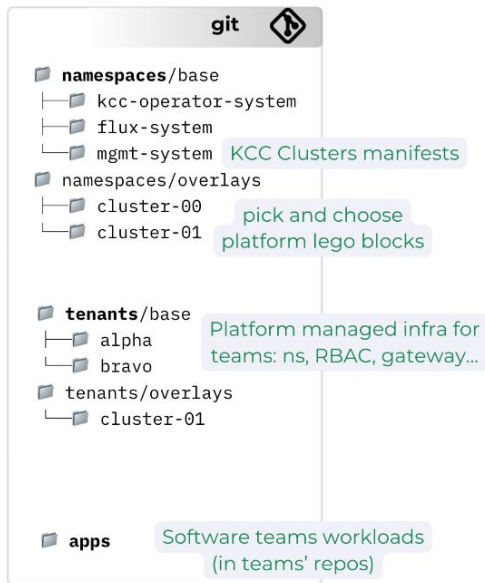
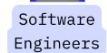
Platform Engineer

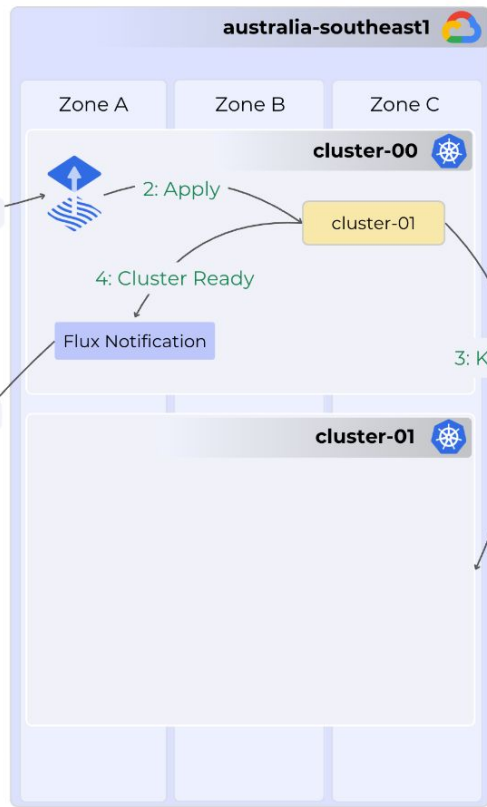

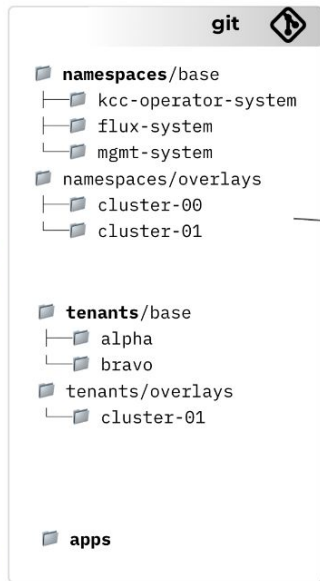
- Cookie-cutter reproducible clusters
- Clusters on-demand: Testing Environments, Disaster Recovery.
- Golden paths and guardrails
- Scalable



Application Developer

- Infrastructure management decoupled from Application
- App placement based on archetype
- Scale seamlessly to 5 or 500 clusters





KCC Custom Resources

1: Sync

2: Apply

4: Cluster Ready

5: Trigger

3: KCC Create

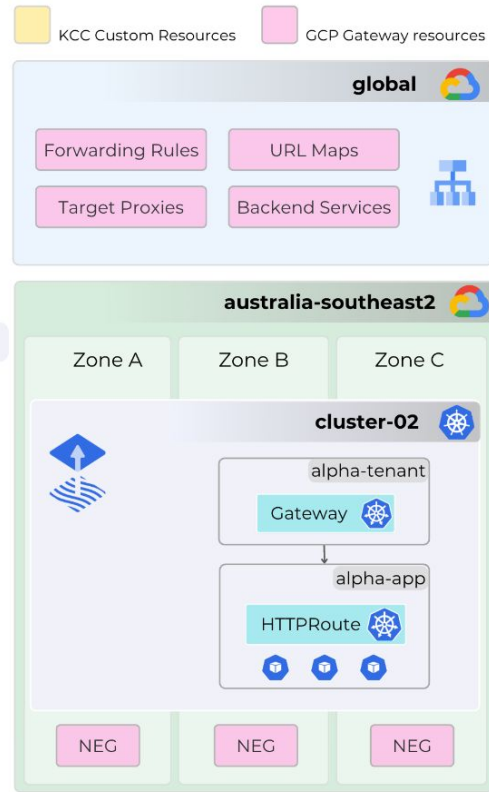
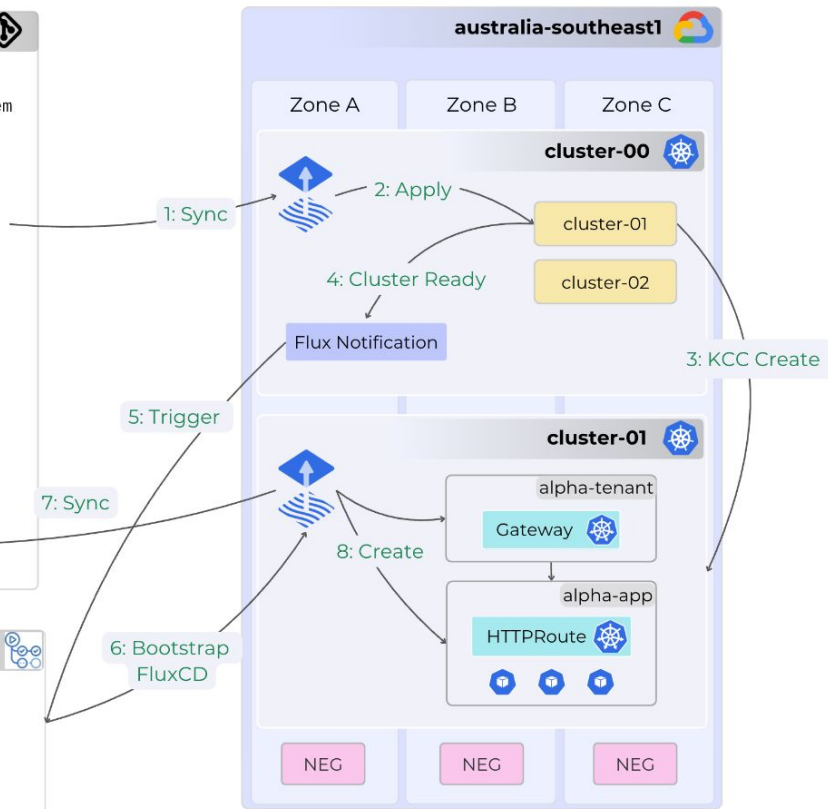
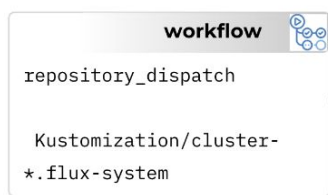
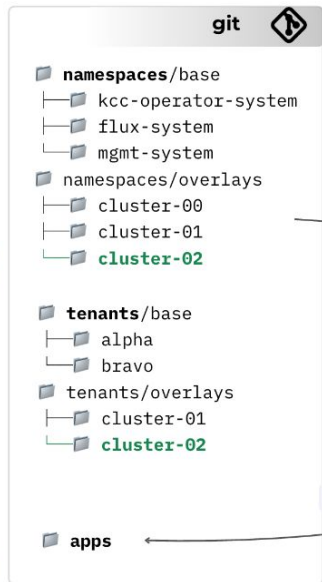
Flux Notification



Platform Engineers



Software Engineers



Takeaways

Unlock the power
of archetypes



Unlock the
power GitOps



Scale kubernetes
configuration



Conclusion & Links

- <https://fluxcd.io/flux/>
- <https://control-plane.io/posts/d2-reference-architecture-guide/>
- <https://rawkode.academy/read/fluxcd-the-inevitable-choice/>
- <https://slsa.dev/spec/v1.2-rc1/use-cases>
- <https://cloud.google.com/architecture/deployment-archetypes>
- <https://opengitops.dev/>
- <https://platformengineering.org/blog/what-is-platform-engineering>


Let's Connect!



[linkedin.com/in/olgamirensky/](https://www.linkedin.com/in/olgamirensky/)



[linkedin.com/in/michael-fornaro/](https://www.linkedin.com/in/michael-fornaro/)



Easygo Careers